

Homework Assignment 5: Heat Diffusion Using Multi-GPU CUDA Programming

Haohan Tsao B09607009

1 Introduction

This report presents the implementation and analysis of solving the heat diffusion equation on a two-dimensional Cartesian grid using CUDA multi-GPU parallel programming. The heat diffusion equation is a fundamental partial differential equation that describes the distribution of heat (or variation of temperature) in a medium over time. In the steady-state condition, this reduces to Laplace's equation.

The steady-state heat diffusion equation is given by:

$$\nabla^2 u = 0 \tag{1}$$

Where:

- u represents the temperature distribution
- ∇^2 is the Laplacian operator

The assignment required solving this equation on a 1024×1024 Cartesian grid with specific boundary conditions:

- Top edge: 400 K (heated boundary)
- Other three edges: 273 K (cooled boundaries)

We were asked to investigate the optimal block size for this problem and compare the performance between single GPU and dual GPU implementations.

2 Theoretical Background

2.1 Heat Diffusion to Laplace Equation

The time-dependent heat diffusion equation is:

$$\frac{\partial u}{\partial t} = \alpha \nabla^2 u \quad (2)$$

Where α is the thermal diffusivity. In steady-state conditions ($\frac{\partial u}{\partial t} = 0$), this becomes Laplace's equation. However, we can use the time-dependent formulation as an iterative method to reach the steady-state solution.

2.2 Discretization on a 2D Grid

On a two-dimensional grid with spacing $\Delta x = \Delta y = 1$, the Laplacian operator can be discretized using the finite difference method:

$$\nabla^2 u(x, y) \approx u(x+1, y) + u(x-1, y) + u(x, y+1) + u(x, y-1) - 4u(x, y) \quad (3)$$

2.3 Iterative Solution Method

For the steady-state solution, we use the iterative Jacobi method:

$$u^{n+1}(x, y) = \frac{1}{4} [u^n(x+1, y) + u^n(x-1, y) + u^n(x, y+1) + u^n(x, y-1)] \quad (4)$$

This iterative approach converges to the solution of Laplace's equation with the given boundary conditions.

3 Experimental Setup

3.1 Implementation Approach

The implementation utilizes CUDA multi-GPU programming with the following key features:

1. **Multi-GPU Architecture:** Using OpenMP to manage multiple GPUs, with one CPU thread per GPU
2. **Domain Decomposition:** The 2D grid is partitioned among GPUs for parallel processing

3. **P2P Communication:** GPU-to-GPU communication for boundary data exchange
4. **Double Buffering:** Two arrays are used to avoid memory copying during iterations

3.2 Testing Methodology

To determine the optimal configuration, we performed experiments with:

GPU Configurations:

- Single GPU (1×1 arrangement)
- Dual GPU (2×1 arrangement)

Block Sizes:

- 8×8 threads per block
- 16×16 threads per block
- 32×32 threads per block

Problem Size:

- Grid size: 1024×1024
- Convergence criterion: $\epsilon = 10^{-10}$
- Maximum iterations: 10,000,000

4 Implementation Details

4.1 Source Code Structure

The implementation consists of the following key components:

- `heat_diffusion.cu`: Main CUDA source file containing host code and device kernels
- `generate_experiments.sh`: Automated experiment generation script
- `submit_all.sh`: Batch job submission script
- `collect_results.sh`: Results collection and analysis script
- `Makefile`: Compilation configuration

4.2 Key CUDA Implementation Features

```
1 __global__ void heat_diffusion(float* phi0_old, float*
  phiL_old,
2                                float* phiR_old, float*
  phiB_old,
3                                float* phiT_old, float*
  phi0_new,
4                                float* C)
5 {
6     // Calculate thread indices
7     int x = blockDim.x*blockIdx.x + threadIdx.x;
8     int y = blockDim.y*blockIdx.y + threadIdx.y;
9     int site = x + y*Lx;
10
11     // Skip boundary points
12     if (skip == 0) {
13         // Apply 4-point stencil
14         phi0_new[site] = 0.25*(b+l+r+t);
15         diff = phi0_new[site] - phi0_old[site];
16     }
17
18     // Parallel reduction for error calculation
19     cache[cacheIndex] = diff*diff;
20     __syncthreads();
21     // ... reduction logic
22 }
```

Listing 1: Core heat diffusion kernel implementation

4.3 Multi-GPU Communication Strategy

For multi-GPU implementation, P2P communication is established between neighboring GPUs:

```
1 // Enable P2P access between neighboring GPUs
2 if (NGPU > 1) {
3     int cpuid_r = ((cpuid_x+1)%NGx) + cpuid_y*NGx;
4     cudaDeviceEnablePeerAccess(Dev[cpuid_r], 0);
5     int cpuid_l = ((cpuid_x+NGx-1)%NGx) + cpuid_y*NGx;
6     cudaDeviceEnablePeerAccess(Dev[cpuid_l], 0);
7 }
```

Listing 2: P2P communication setup

5 Results and Analysis

5.1 Performance Comparison

The following table summarizes the performance results for all tested configurations:

GPU Count	Block Size	Time (ms)	Gflops	Iterations
1	8×8	124,809	43.13	736,246
1	16×16	96,977	55.51	736,246
1	32×32	120,626	44.63	736,246
2	8×8	80,869	66.56	736,246
2	16×16	57,593	93.47	736,246
2	32×32	66,007	81.55	736,246

Table 1: Performance comparison for different GPU and block size configurations

5.2 Optimal Block Size Analysis

The results clearly demonstrate that **16×16 block size is optimal** for this problem:

Block Size	Single GPU Gflops	Dual GPU Gflops	Performance Ratio
8×8	43.13	66.56	1.54×
16×16	55.51	93.47	1.68×
32×32	44.63	81.55	1.83×

Table 2: Performance comparison across block sizes

5.3 Multi-GPU Scaling Analysis

The multi-GPU implementation shows excellent scaling efficiency:

- **Best Configuration:** 2 GPUs with 16×16 blocks achieving 93.47 Gflops
- **Speedup:** 1.68× improvement over single GPU
- **Efficiency:** 84% of ideal 2× speedup
- **Consistent Convergence:** All configurations converged in exactly 736,246 iterations

6 Discussion

6.1 Block Size Optimization

The superior performance of 16×16 blocks can be attributed to several factors:

- **Memory Coalescing:** 16×16 provides optimal memory access patterns for the 2D grid
- **Occupancy:** Balances thread count (256 threads) with resource utilization
- **Shared Memory Usage:** Efficient utilization without bank conflicts
- **Warp Alignment:** Better alignment with GPU warp architecture (32 threads)

The poor performance of 32×32 blocks (1024 threads) suggests resource limitations, while 8×8 blocks underutilize the GPU cores.

6.2 Multi-GPU Efficiency

The multi-GPU implementation demonstrates strong scaling:

- **Load Balancing:** Even distribution of computational work
- **Communication Overhead:** Minimal impact from P2P boundary exchange
- **Synchronization:** Efficient OpenMP barriers between iterations
- **Memory Bandwidth:** Effective utilization of aggregate memory bandwidth

6.3 Convergence Behavior

All configurations converged to the same solution in identical iteration counts, confirming:

- **Mathematical Consistency:** GPU count and block size do not affect convergence
- **Numerical Accuracy:** Proper implementation of boundary conditions
- **Algorithm Stability:** Robust iterative solver implementation

7 Conclusion

This study successfully demonstrates the implementation and optimization of a multi-GPU CUDA solver for the heat diffusion equation. The key findings are:

1. **Optimal Configuration:** 2 GPUs with 16×16 block size provides the best performance at 93.47 Gflops
2. **Block Size Impact:** 16×16 consistently outperforms other block sizes across all GPU configurations
3. **Multi-GPU Scaling:** Achieves 84% efficiency with $1.68 \times$ speedup using dual GPUs
4. **Robust Implementation:** Consistent convergence behavior regardless of parallelization strategy

The implementation showcases the effectiveness of multi-GPU computing for large-scale numerical simulations, with proper domain decomposition and P2P communication enabling near-linear scaling. The 16×16 block size emerges as the optimal choice, balancing computational efficiency with hardware resource utilization.

For production use, the dual-GPU configuration with 16×16 blocks is recommended for maximum performance, while the single-GPU implementation with 16×16 blocks provides an efficient solution for memory-constrained scenarios.

References

- CUDA Programming Guide, NVIDIA Corporation
- OpenMP Specification, OpenMP Architecture Review Board
- Numerical Methods for Partial Differential Equations, Morton & Mayers